

MSiA-413 Introduction to Databases and Information Retrieval

Lecture 2 Fixed-point and Floating-point Representations

Instructor: Nikos Hardavellas

Slides adapted from Steve Tarzia

Last week we talked about Integers

- Integers can be stored in a base-two positional notation in binary
- Addition and subtraction follow the familiar mechanics
- Learned some tricks (eg., $2^{10} = 1024 \approx 1000$)
- Signed integers use two's complement representation for negatives
 - Two's complement makes subtraction just as easy as addition
 - Positive numbers are represented in the same way whether you're using a signed or unsigned data type, but
 - Small negatives and huge positives can be confused if you misinterpret the type

A few more things about integers

- Multiplication: two's complement works magically here too
- Positive division works as expected
- “*Sign extension*.” when increasing the “bit size” of a negative number, add leading ones
 - Eg., -2 is **1110** as a 4-bit signed integer and **11111110** in 8 bits
- Computers typically use 32 or 64 bit integers


Any questions on last week's material?

3

Integers are great for **counting**, but sometimes we need to **measure** fractional quantities

Binary numbers can have “decimal” places, too

- **0.11111111**_{two} is slightly smaller than 1
- **0.000000001**_{two} is slightly larger than 0
- **0.1**_{two} is one half


$$\begin{aligned} \bullet \text{ } 10.101_{\text{two}} &= \textcolor{red}{1} \times 2^1 + \textcolor{red}{0} \times 2^0 + \textcolor{red}{1} \times 2^{-1} + \textcolor{red}{0} \times 2^{-2} + \textcolor{red}{1} \times 2^{-3} \\ &= 2 + 0 + 1/2 + 0 + 1/8 = 2 = 2.625_{\text{ten}} \end{aligned}$$

How shall we represent fractional number in the computer?


4

Fixed point: Integers 2.0

- Simplest solution is to just stick an implicit **binary (radix) point** somewhere (We don't call it a decimal point because we're not in base ten)
- Examples of fixed point numbers in base ten:
 - Represent the cost of a purchase with an integer number of cents
 - The cost of a sandwich is 625 cents (\$6.25)
 - Represent the distance between cities by counting the hundredths of a mile
 - Evanston is 1321 hundredths of a mile from Chicago (13.21 miles)
 - and 79,543 hundredths of a mile from Philadelphia

5

Fixed point example in 16 bits



Let's store the chemical elements' atomic weights

- Smallest value (hydrogen) is 1.00784
- Largest value (uranium) is 238.02891
- Negative values are not possible
- We can reserve 8 bits for the fractional part and 8 bits for the part > 1
- In this particular binary fixed point representation, the weight of uranium is:
 1110111000000111_2 Remember that the radix point is implicit. This represents the value
 $11101110.00000111_2 = 238_{10} = 238.02734375_{10}$
(We had to round off, so this is not precisely accurate)
- And the weight of hydrogen is:
 0000000100000010_2 ,
i.e., $00000001.00000010 = 1 = 1.0078125$

6

Fixed point is simple & efficient but it has its limitations

- Range is very limited
 - Multiplication overflows easily – can double the number of bits
 - E.g., if working in 32-bits, then we can only multiply 16-bit values without overflow
 - Division **underflows** easily (small values are rounded to zero)
- Precision varies across the range:
 - Small numbers have few significant figures

7


Floating point

- Based on scientific notation:
 - $10,340 = 1.034 \times 10^4$
 - $0.00424 = 4.24 \times 10^{-3}$
- Gives a compact representation of extreme values:
 - $1,000,000,000,000,000,000,000 = 1.0 \times 10^{24}$
 - $0.000\ 000\ 000\ 000\ 000\ 000\ 000\ 001 = 1.0 \times 10^{-24}$
- In binary:
 - $100010_{\text{two}} = 1.0001_{\text{two}} \times 2^5_{\text{ten}} = 1.0001 \times 10^{101}_{\text{two}}$
 - $0.00101_{\text{two}} = 1.01_{\text{two}} \times 2^{-3}_{\text{ten}} = 1.01 \times 10^{-11}_{\text{two}}$

8

Representing floating point in bits


- $0.15625_{\text{ten}} = 0.00101_{\text{two}} = 1.01_{\text{two}} \times 2^{-3}_{\text{ten}} = 1.01 \times 10^{-11}_{\text{two}}$
- Three essential parts are the **sign**, **fraction**, & **exponent**
 - Notice that the first significant figure is always “1” so we don’t have to store it
- In the mid 1980s, the IEEE standardized the floating point representation of 32 and 64 bit numbers:
 - The exponent has a sign too, but the standard says to add a “bias” of 127



9

64-bit floating point

- Similar to 32-bit, but we have more precision in the fraction and larger exponents are possible
- 32-bit is called **single precision** and 64-bit is called **double precision**
- Double precision can represent larger, smaller, and more precise numbers



10

A few special floats


- The IEEE standard allows for a few special values to be stored
 - Positive and negative zero (remember that we normally start with an implied “1”)
 - All exponent bits set to zeros
 - Positive and negative infinity (the result of divide by zero)
 - Not a number (the result of zero divided by zero)
 - These all have the exponent bits set to all ones

The Flexibility and Flaws of Floats

- A 32-bit signed integer can represent all the whole numbers between -2,147,483,648 and 2,147,483,647
- A 32-bit floating point number can be as large as $\pm 3.402823 \times 10^{38}$
 $= 340,282,300,000,000,000,000,000,000,000,000,000,000,000$
- or as tiny as $5.8774718 \times 10^{-39}$
 $= 0.000\ 000\ 000\ 000\ 000\ 000\ 000\ 000\ 000\ 000\ 000\ 000\ 005\ 877\ 471\ 8$
- But, single-precision floats have only 24 bits of precision:
 - Can only precisely store integers up to $2^{24} = 16,777,216$
- Floats can store larger numbers than integers of the same bit-length, but with less precision because 8 bits are set aside for the exponent

12

Floats just distributed the same number of values differently – with exponential spacing



13

Know when to use integers, floating point, and fixed point

- When **counting** or labelling things, always use integers
- When **measuring** physical quantities, usually use floating point
 - May use fixed point if speed/simplicity is more important than accuracy
- If your machine does not support floating point (eg., a toaster):
 - Use fixed point representation for fractional quantities
- If rounding is desired then use fixed point (but carefully)
 - U.S. currency values usually should be rounded to the nearest cent
- Use 64-bit integers when you need values > 2 billion
- Floating point rules of thumb:
 - Single precision gives ~ 7 decimal digits of precision
 - Double precision gives ~ 16 decimal digits of precision

14

One more point about fractions in binary: Base ten decimals usually have to be rounded

- We all know that $1/3$ cannot be represented exactly in decimal
 - That's because 10^x not divisible by 3 (for any integer x)
- Similarly, $1/10$ cannot be represented exactly in binary
 - Because 2^x is not divisible by 10 (for any integer x)
- In general, a rational number a/b can be *exactly* represented in binary only if b is a power of 2
 - Otherwise, there is some rounding error
- Most fractions cannot be stored exactly with a finite number of bits
 - Actually, this is also true in decimal!
- So, always expect small rounding errors when working in floating point

15

How do computers work with floats?

- It's complicated and slow!
- Have to manipulate both the fraction and the exponent
- Addition is no longer simple

16