

# MSiA-413 Introduction to Databases and Information Retrieval

## Lecture 16 Triggers, Introduction to Transactions

Instructor: Nikos Hardavellas

Slides adapted from S. Tarzia, A. Silberschatz, H.F. Korth, S. Sudarshan

## Last Lecture

- **Recursive Queries on Networks**
  - Powerful queries to express arbitrarily long chains of dependent queries
- **Views**
  - Temporary relations that are not part of the conceptual model
- **WITH** statement
  - Creates scoped views available only to the query issuing the WITH clause
- Set comparison and existential operators
  - **SOME, ANY, ALL, EXISTS**

## Part I

### Triggers

3

## Adding integrity constraints

SalesOrders.sqlite

- *New rule: no more than 4 employees in the same office (i.e., with same area code)*

```
SELECT EmpAreaCode, COUNT(*) AS NumEmployeesAtOffice
FROM Employees
GROUP BY EmpAreaCode;
```

EmpAreaCode	NumEmployeesAtOffice
206	1
210	1
253	2
425	4
515	1

- Inserting new employee at area code 425 should fail

```
INSERT INTO Employees
VALUES (800, "Nikos", "Hardavellas", "2233 Tech Dr",
      "Evanston", "IL", "60208", 425, "491-2270");
```

4

## How to enforce the rule? Triggers

- A **trigger** is a statement that is executed automatically by the system as a side effect of a modification to the database
- To design a trigger mechanism, we must:
  - Specify the conditions under which the trigger is to be executed.
  - Specify the actions to be taken when the trigger executes

specify when and what to do

5

## Adding a trigger

SalesOrders.sqlite

- New rule: no more than 4 employees in the same office (i.e., with same area code)

```
CREATE TRIGGER Max4EmployeesPerOffice
BEFORE INSERT
ON Employees
FOR EACH ROW
BEGIN
    SELECT CASE
        WHEN (SELECT COUNT(*)
              FROM Employees
              WHERE EmpAreaCode = new.EmpAreaCode) >= 4
        THEN RAISE(FAIL, "Error: max 4 employees per office")
        END;
END;
```

*Trigger actions* {

Trigger name

When to "fire"

On which table

Run trigger per row

Refer to "new row"

new. = the row that the insert statement attempts to insert

Raise an exception & print error

6

## Insert with the trigger defined

SalesOrders.sqlite

- Inserting new employee at area code 425 should fail

```
INSERT INTO Employees  
VALUES (800, "Nikos", "Hardavellas", "2233 Tech Dr",  
       "Evanston", "IL", "60208", 425, "491-2270");
```

Error: max 4 employees per office: INSERT INTO Employees  
VALUES (800, "Nikos", "Hardavellas", "2233 Tech Dr",  
 "Evanston", "IL", "60208", 425, "491-2270")

- Triggers are supported by SQLite, **but not** DB Browser for SQLite
  - trigger bug report, github: *"We practically don't handle triggers at all in our application because they are complicated to parse"*
- Use the sqlite3 command-line interface if in trouble

7

## Trigger Events

- The triggering event can be **insert**, **delete**, **update**, or **update of**

```
CREATE TRIGGER trigger_name  
BEFORE INSERT ON table  
...
```

```
CREATE TRIGGER trigger_name  
BEFORE DELETE ON table  
...
```

```
CREATE TRIGGER trigger_name  
BEFORE UPDATE ON table  
...
```

```
CREATE TRIGGER trigger_name  
BEFORE UPDATE OF column1, column2, ... ON table  
...
```

8

## Trigger Timing

- The trigger can fire **before**, **after**, or **instead of** the triggering event

```
CREATE TRIGGER trigger_name  
BEFORE INSERT ON table  
...
```

← Typical use: add  
integrity constraints

```
CREATE TRIGGER trigger_name  
INSERT ON table  
...
```

← Default is BEFORE

```
CREATE TRIGGER trigger_name  
AFTER INSERT ON table  
...
```

← Typical use: perform  
additional actions

```
CREATE TRIGGER trigger_name  
INSTEAD OF INSERT ON table  
...
```

9

## INSTEAD OF Triggers

take another action than the user specified action

- Changes the statement to execute
- Example: instead of modifying a view, modify the main table

```
CREATE TABLE Customers(ID INTEGER PRIMARY KEY,  
                        Name varchar(100),  
                        Addr varchar(100) );
```

```
CREATE VIEW CustomerView  
AS SELECT ID, Addr FROM Customers;
```

```
CREATE TRIGGER CustomerViewChange  
INSTEAD OF UPDATE OF Addr ON CustomerView  
BEGIN  
    UPDATE customers SET Addr = new.Addr  
    WHERE ID = new.ID;  
END;
```

10

## Referencing attributes of old/new rows

- Use **old.** and **new.** to refer to the old/new rows of the insert, update, or delete statement that fired the trigger
- **INSERT**: **new.** references are valid **Reference all the rows with new values**
- **UPDATE**: **new.** and **old.** references are valid
- **DELETE**: **old.** references are valid **Reference all the rows with original values**
- Example:

```
CREATE TRIGGER AddrChange
AFTER UPDATE OF Addr ON Customers
BEGIN
    UPDATE orders SET Addr = new.Addr
    WHERE ID = old.ID;
END;
```

Refers to the row  
after the update

Refers to the row  
before the update

11

## Trigger execution granularity

- Defines how often the trigger will execute
- Example:

```
CREATE TRIGGER TriggerName
BEFORE INSERT
ON table
[FOR EACH ROW | FOR EACH STATEMENT]
BEGIN
    ...
END;
```

Execute trigger once for  
each row inserted by the  
triggering statement

Execute trigger once for  
the triggering statement

- SQLite implements only per-row triggers, hence this clause is optional

12

## Trigger event filtering

- Execute trigger only when certain conditions are satisfied
- Example:

```
CREATE TRIGGER TriggerName
BEFORE INSERT
ON table
WHEN condition      similar to WHERE clause
BEGIN
...
END;
```

- SQLite implements only per-row triggers, hence this clause is optional

13

## Raising exceptions

- Notify the caller that an error has occurred
  - Actions: print an error message, return an error to the application if needed
- **RAISE** is an **expression**, not a statement
  - Must be within a **SELECT**, **CASE**, or any other statement accepting expressions

```
CREATE TRIGGER Max4EmployeesPerOffice
BEFORE INSERT ON Employees
FOR EACH ROW
BEGIN
  SELECT CASE
    WHEN (SELECT COUNT(*)
          FROM Employees
          WHERE EmpAreaCode = new.EmpAreaCode) >= 4
    THEN RAISE(FAIL, "Error: max 4 employees per office")
  END;
END;
```

Conflict resolution algorithm

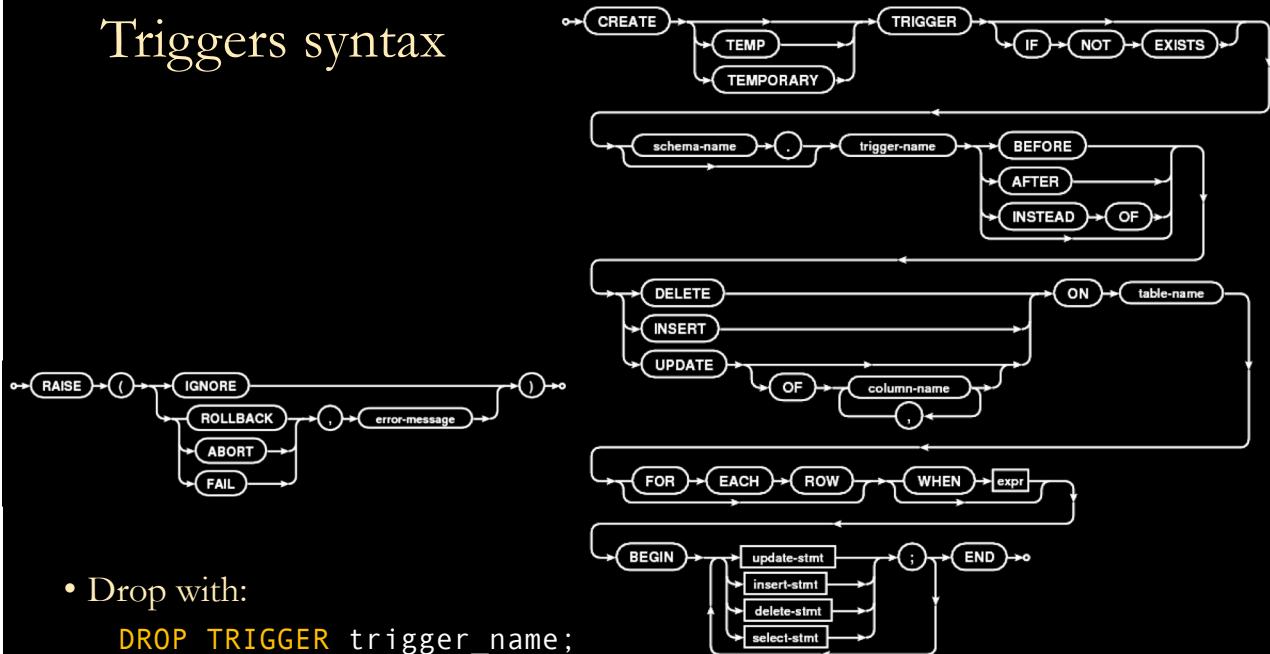
14

# Conflict resolution algorithms

- **FAIL**: stop processing the rest of the current SQL statement
  - Do not undo any prior changes don't undo what didn't fail
  - If it fails on the 100<sup>th</sup> row, the actions taken due to the previous 99 rows are preserved
  - The transaction remains active (if within one)
- **ABORT**: stop processing the rest of the current SQL statement and abort
  - **Undo** any prior changes made by the **current SQL statement**
  - Changes caused by prior SQL statements within the same transaction are preserved
  - The transaction remains active
- **ROLLBACK**: stop processing the rest of the current SQL statement and rollback
  - **Undo** any prior changes made by **all SQL statements** in the transaction
  - **End the transaction**
  - If not within a transaction, ROLLBACK and ABORT are the same
- **IGNORE**: skip the one row that violates the constraint skip the one that violates
  - Continue processing subsequent rows as if nothing went wrong
  - Do not return an error to the application

15

# Triggers syntax



16



## Part II

### Transactions

17

## Transaction Concept

unit of work - single atomic element - happens or not

- A **transaction** is a unit of program execution that accesses and possibly updates various data items
- Example: transaction to transfer \$50 from account A to account B:
  1. read(A)
  2.  $A := A - 50$
  3. write(A)
  4. read(B)
  5.  $B := B + 50$
  6. write(B)
- Two main issues to deal with:
  - Failures of various kinds, such as hardware failures and system crashes
  - Concurrent execution of multiple transactions

18

## Atomicity requirement

- Transaction to transfer \$50 from account A to account B:
  1. `read(A)`
  2. `A := A - 50`
  3. `write(A)`    transfer needs to happen in entirety or not at all
  4. `read(B)`
  5. `B := B + 50`
  6. `write(B)`
- What if the transaction fails at step 5 ?
  - Money will be “lost” leading to an inconsistent database state
  - Failure could be due to software or hardware
- The system should ensure that updates of a partially executed transaction are not reflected in the database

19

## Durability requirement

- Transaction to transfer \$50 from account A to account B:
  1. `read(A)`
  2. `A := A - 50`
  3. `write(A)`
  4. `read(B)`
  5. `B := B + 50`
  6. `write(B)`
- The updates to the database by the transaction must persist even if there are software or hardware failures
- Once the user has been notified that the transaction has completed (i.e., the transfer of the \$50 has taken place) the state of the database should always reflect that

20

## Consistency requirement

- Transaction to transfer \$50 from account A to account B:
  1. `read(A)`
  2. `A := A - 50`
  3. `write(A)`
  4. `read(B)`
  5. `B := B + 50`
  6. `write(B)`

check sum integrity constraint  
before and after transaction
- In above example: the sum of A and B is unchanged
- In general, consistency requirements include
  - Explicit integrity constraints, e.g., primary keys, foreign keys, unique values
  - Implicit integrity constraints, e.g., balances minus loans must equal cash-in-hand
- A transaction must see a consistent database
  - During transaction execution the database may be temporarily inconsistent
  - When the transaction completes successfully the database must be consistent

21

## Isolation requirement

- Transaction to transfer \$50 from account A to account B:

User 1	User 2
1. <code>read(A)</code>	
2. <code>A := A - 50</code>	
3. <code>write(A)</code>	
	<code>read(A), read(B), print(A+B)</code>
4. <code>read(B)</code>	
5. <code>B := B + 50</code>	
6. <code>write(B)</code>	
- User 2 should not be allowed to see the temporarily inconsistent database
  - The sum A+B should not be incorrect, otherwise money appear to be “lost”
- Provide the illusion that transactions execute **serially**, i.e., one after the other
  - User 1 fully executes his transaction, then User 2 fully executes his transaction
  - ...or the other way around

22

## ACID properties

- **Atomicity.** Either all operations of the transaction are properly reflected in the database, or none are
- **Consistency.** The execution of a transaction in isolation preserves the consistency of the database
- **Isolation.** Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions
  - Intermediate results must be hidden from the outside world
  - For every pair of transactions  $T_i$  and  $T_j$ , it appears to  $T_i$  that
    - Either  $T_i$  finished execution before  $T_j$  started, or
    - $T_j$  finished execution before  $T_i$  started
- **Durability.** After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures

23

## Transaction state

- **Active:** the initial state; the transaction stays in this state while it is executing
- **Partially committed:** after the final statement has been executed
- **Failed:** after the discovery that normal execution can no longer proceed **temporary state**
- **Aborted:** after the transaction has been rolled back and the database restored to its state prior to the start of the transaction.
  - Two options after it has been aborted:
    - Restart the transaction (can be done only if no internal logical error )
    - Kill the transaction
- **Committed:** after successful completion
- Transactions begin implicitly or explicitly
  - Ended by commit work or rollback work
- Default on most databases: each SQL statement commits automatically

24

## Example: atomically delete rows

SalesOrders.sqlite

```
select * from order_details;
```

Commit transaction

Begin transaction

```
savepoint transaction1;  
delete from order_details where orderNumber=1;  
delete from order_details where orderNumber=2;  
select * from order_details;  
commit;  
  
select * from order_details;
```

	OrderNumber	ProductNumber	QuotedPrice	QuantityOrdered
1	1	1	1200	2
2	1	6	635	3
3	1	11	1650	4
4	1	16	28	1
5	1	21	55	3
6	1	26	121.25	5
7	1	40	174.6	6
8	2	27	24	4
9	2	40	180	4
10	3	1	1164	5

	OrderNumber	ProductNumber	QuotedPrice	QuantityOrdered
1	3	1	1164	5
2	3	6	615.95	5
3	3	11	1650	1
4	3	16	28	2

25

## Example: rollback attempt to atomically delete rows

SalesOrders.sqlite

```
select * from order_details;
```

Rollback and end transaction

Begin transaction

```
savepoint transaction2;  
delete from order_details where orderNumber=3;  
delete from order_details where orderNumber=4;  
select * from order_details;  
rollback;  
  
select * from order_details;
```

	OrderNumber	ProductNumber	QuotedPrice	QuantityOrdered
1	3	1	1164	5
2	3	6	615.95	5
3	3	11	1650	1
4	3	16	28	2

	OrderNumber	ProductNumber	QuotedPrice	QuantityOrdered
1	3	1	1164	5
2	3	6	615.95	5
3	3	11	1650	1
4	3	16	28	2

26

# Named transactions

- **BEGIN ... END** is another way to denote a transaction
  - **END** and **COMMIT** are the same: complete and exit the transaction
  - **ROLLBACK**: undo all changes and cancel the transaction
    - Subsequent SQL statements are not part of a transaction
  - **BEGIN** cannot be used within a transaction (i.e., no nesting)
- **SAVEPOINT** starts a transaction that is named and can be nested
  - **COMMIT**: commits all outstanding transactions and leaves transaction stack empty
  - **ROLLBACK**: undo all changes and cancel the transaction
  - **ROLLBACK TO TransactionName**:
    - Undo all changes until the beginning of TransactionName, and
    - Enter a new transaction with the name TransactionName
  - **RELEASE TransactionName**:
    - Remove all savepoints back to and including the savepoint named TransactionName
    - Cannot rollback to these savepoints anymore
    - No write back of modifications; **COMMIT** does that

27

## Example: nested named transactions

SalesOrders.sqlite

```
select * from order_details;
```

	OrderNumber	ProductNumber	QuotedPrice	QuantityOrdered
1	3	1	1164	5
2	3	6	615.95	5
3	3	11	1650	1
4	3	16	28	2

```
T1: savepoint transaction1;
```

```
T1: delete from order_details where orderNumber=3;
```

```
T1: savepoint transaction2;
```

```
T2: delete from order_details where orderNumber=4;
```

```
T2: select * from order_details;
```

	OrderNumber	ProductNumber	QuotedPrice	QuantityOrdered
1	3	1	1164	5
2	3	6	615.95	5
3	3	11	1650	1
4	3	16	28	2

```
T2: rollback to transaction2;
```

```
T2: select * from order_details;
```

```
T1: still not terminated
```

```
T2: rollback to transaction1;
```

```
T1: rollback;
```

```
select * from order_details;
```

28

release transaction 2